

CPSC 532W Homework 4

Naomi Graham

March 9, 2021

All the code can be found on: https://github.com/n6graham/cpsc532_hw4.

1 BBVI code

For BBVI, I used the graph-based implementation based off of Jason's HW2 code.

1.1 BBVI Algorithm

Even though I was getting good proposals from running SGD, it seems that a bug in my importance sampling was causing it to look like I had really bad values after sampling.

For example, In program 1, my variational posterior ended up being:

```
1 sigma['Q'] is {'sample2': Normal(loc: 8.311932563781738, scale :  
2 1.0169004201889038)}
```

But the posterior mean and variance returned by importance sampling were:
tensor(1.7418) Variance: tensor(2.1257)

1.2 Code

I implemented SGD for optimizer-step. Also returned the value max-norm to keep track of how the norm of the gradient was looking. It was almost always somewhat-monotonically decreasing.

```
1 def optimizer_step(q, ghat, t):  
2     for v, d in q.items():  
3         print("norm of gradient:", np.linalg.norm(ghat[v]))  
4         i = 0  
5         for params in d.Parameters():  
6             params.data = params.data + ghat[v][i]/(t+10)  
7  
8     max_norm = np.max([ np.linalg.norm(ghat[v]) for v, d in q.items() ])  
9  
10    return q, max_norm  
11
```

I also implemented elbo-grad, which was a lot of work (but worth it, of course!)

```

1  def elbo_grad(Glist, logWlist):
2      L = len(Glist)
3      Flist = list([{} for i in range(0,L)])
4      ghat = {}
5      U = list(set([u for G in Glist for u in G]))
6      print("here!!!")
7
8      for v in U:
9          # get number of parameters for v
10         for i in range(0,L):
11             if v in list(Glist[i].keys()):
12                 num_params = len(Glist[i][v])
13                 break
14
15         for i in range(0,L):
16             if v in list(Glist[i].keys()):
17                 x = Glist[i][v]*logWlist[i]
18                 if i == (L-1): print("x is ", x)
19                 Flist[i][v] = x
20             else:
21                 Flist[i][v] = torch.tensor([0 for j in range(
num_params)])
22                 Glist[i][v] = torch.tensor([0 for j in range(
num_params)])
23
24         Fv = [ Flist[i][v] for i in range(0,L)]
25         Gv = [ Glist[i][v] for i in range(0,L)]
26         Fv = torch.stack(Fv)
27         Gv = torch.stack(Gv)
28         Fv = Fv.detach().numpy()
29
30         varG = [ np.var(np.array(Gv[:,j])) for j in range(num_params)
]
31         denom = sum(varG)
32
33         C = np.array([ np.cov(Fv[:,j],Gv[:,j], rowvar=True) for j in
range(num_params) ])
34         cov = [ C[j][1][0] for j in range(num_params) ]
35         numerator = sum(cov)
36         bhat = numerator/denom
37
38         print("bhat is", bhat)
39
40         #numerator = np.array([ np.sum(C[j]) for j in range(num_params
) ])
41
42
43         ghat[v] = sum( np.divide((Fv - bhat*np.array(Gv)),L) )
44
45
46         print("returning ghat:", ghat)
47
48         return ghat
49
50

```

Here I implement algorithm 15.

```

1      sigma = {'Q':{}, 'logW':0, 'G':{}}
2      weighted_samples = []
3
4      for t in range(0,T): # T is the number of iterations
5          Glist = []
6          logWlist = []
7
8          # here we compute a batch of gradients
9          for l in range(0,L): # L is the batch size
10             sigma['logW']=0
11             # first we get the trace and update sigma using sample from
joint
12             #val_l, sigma_l, trace_l = sample_from_joint(graph,sigma)
13             r_l, sigma_l, trace_l = sample_from_joint(graph,sigma)
14             # then we get the deterministic expression using the trace
15             #deterministic_expr = plugin_parent_values(expr,trace_l)
16             G_l = copy.deepcopy(sigma_l['G'])
17             logW_l = sigma_l['logW']
18             Glist.append(G_l)
19             logWlist.append(logW_l)
20
21             weighted_samples.append((r_l,logW_l))
22
23             ELBO = sum(logWlist)
24
25             ghat = elbo_grad(Glist,logWlist)
26
27             sigma['Q'], max_norm = optimizer_step( sigma['Q'],ghat,t) #update
the proposal
28             print("results on iteration {} are ".format(t), sigma['Q'])
29             print("the max gradient is ", max_norm )
30
31
32
33             print("sigma['Q'] is ",sigma['Q'])
34
35             return weighted_samples, sigma['Q']
36


```

Sadly my importance sampling just was not working for some reason :(

2 Program 1

```
returning ghat: {'sample2': array([-0.49577966, -0.0261793 ], dtype=float32)}
norm of gradient: 0.49647036
results on iteration 93 are {'sample2': Normal(loc: 7.7331366539001465, scale: 0.91798853874
20654)}
the max gradient is 0.49647036
ELBO is tensor(-951.8278, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.5220273 , 0.08419174], dtype=float32)}
norm of gradient: 0.5287729
results on iteration 94 are {'sample2': Normal(loc: 7.728116989135742, scale: 0.917835891246
7957)}
the max gradient is 0.5287729
ELBO is tensor(-968.4595, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.42646474, -0.03271492], dtype=float32)}
norm of gradient: 0.42771772
results on iteration 95 are {'sample2': Normal(loc: 7.724055290222168, scale: 0.918322205543
5181)}
the max gradient is 0.42771772
ELBO is tensor(-951.6724, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.39113718, 0.0641632 ], dtype=float32)}
norm of gradient: 0.39636502
results on iteration 96 are {'sample2': Normal(loc: 7.720365524291992, scale: 0.918134987354
2786)}
the max gradient is 0.39636502
ELBO is tensor(-975.7202, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.42080986, 0.08003196], dtype=float32)}
norm of gradient: 0.4283527
results on iteration 97 are {'sample2': Normal(loc: 7.716432571411133, scale: 0.918498694896
698)}
the max gradient is 0.4283527
ELBO is tensor(-1007.1896, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.38780287, -0.11757746], dtype=float32)}
norm of gradient: 0.40523517
results on iteration 98 are {'sample2': Normal(loc: 7.712841987609863, scale: 0.918948173522
9492)}
the max gradient is 0.40523517
ELBO is tensor(-1002.4915, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.44282976, -0.13143373], dtype=float32)}
norm of gradient: 0.46192318
results on iteration 99 are {'sample2': Normal(loc: 7.708779335021973, scale: 0.918293952941
8945)}
the max gradient is 0.46192318
sigma['Q'] is {'sample2': Normal(loc: 7.708779335021973, scale: 0.9182939529418945)}
{'sample2': Normal(loc: 7.708779335021973, scale: 0.9182939529418945)}
```

But the sampled values look wrong:



3 Program2

4 Program3

5 Program4

6 Program5

7 Program3

7.1 results