

CPSC 532W Homework 4

Naomi Graham

March 10, 2021

All the code can be found on: https://github.com/n6graham/cpsc_hw4.
Sadly, again, I ran out of time for plotting, but I really think my code works well in theory!

1 BBVI code

For BBVI, I used the graph-based implementation based off of Jason's HW2 code.

1.1 BBVI Algorithm

Even though I was getting good proposals from running SGD, it seems that a bug in my importance sampling was causing it to look like I had really bad values after sampling.

For example, In program 1, my variational posterior ended up being:

```
1 sigma['Q'] is {'sample2': Normal(loc: 7.149489402770996, scale :  
2 0.9122980237007141)}
```

But the posterior mean and variance returned by importance sampling were:

Mean: 1.3713 Variance: 3.5675.

1.2 Code

I implemented SGD for optimizer-step. Also returned the value max-norm to keep track of how the norm of the gradient was looking. It was almost always somewhat-monotonically decreasing.

```
1 def optimizer_step(q, ghat, t):  
2     for v, d in q.items():  
3         print("norm of gradient:", np.linalg.norm(ghat[v]))  
4         i = 0  
5         for params in d.Parameters():  
6             params.data = params.data + ghat[v][i]/(t+10)  
7  
8     max_norm = np.max([ np.linalg.norm(ghat[v]) for v, d in q.items() ])  
9  
10    return q, max_norm  
11
```

I also implemented elbo-grad, which was a lot of work (but worth it, of course!)

```

1  def elbo_grad(Glist, logWlist):
2      L = len(Glist)
3      Flist = list([{} for i in range(0,L)])
4      ghat = {}
5      U = list(set([u for G in Glist for u in G]))
6      print("here!!!")
7
8      for v in U:
9          # get number of parameters for v
10         for i in range(0,L):
11             if v in list(Glist[i].keys()):
12                 num_params = len(Glist[i][v])
13                 break
14
15         for i in range(0,L):
16             if v in list(Glist[i].keys()):
17                 x = Glist[i][v]*logWlist[i]
18                 if i == (L-1): print("x is ", x)
19                 Flist[i][v] = x
20             else:
21                 Flist[i][v] = torch.tensor([0 for j in range(
num_params)])
22                 Glist[i][v] = torch.tensor([0 for j in range(
num_params)])
23
24         Fv = [ Flist[i][v] for i in range(0,L)]
25         Gv = [ Glist[i][v] for i in range(0,L)]
26         Fv = torch.stack(Fv)
27         Gv = torch.stack(Gv)
28         Fv = Fv.detach().numpy()
29
30         varG = [ np.var(np.array(Gv[:,j])) for j in range(num_params)
]
31
32         denom = sum(varG)
33
34         C = np.array([ np.cov(Fv[:,j],Gv[:,j], rowvar=True) for j in
range(num_params) ])
35         cov = [ C[j][1][0] for j in range(num_params) ]
36         numerator = sum(cov)
37         bhat = numerator/denom
38
39         print("bhat is", bhat)
40
41         #numerator = np.array([ np.sum(C[j]) for j in range(num_params
) ])
42
43
44         ghat[v] = sum( np.divide((Fv - bhat*np.array(Gv)),L) )
45
46
47         print("returning ghat:", ghat)
48
49         return ghat
50

```

Here I implement algorithm 15.

```

1      sigma = {'Q':{}, 'logW':0, 'G':{}}
2      weighted_samples = []
3
4      for t in range(0,T): # T is the number of iterations
5          Glist = []
6          logWlist = []
7
8          # here we compute a batch of gradients
9          for l in range(0,L): # L is the batch size
10             sigma['logW']=0
11             # first we get the trace and update sigma using sample from
joint
12             #val_l, sigma_l, trace_l = sample_from_joint(graph,sigma)
13             r_l, sigma_l, trace_l = sample_from_joint(graph,sigma)
14             # then we get the deterministic expression using the trace
15             #deterministic_expr = plugin_parent_values(expr,trace_l)
16             G_l = copy.deepcopy(sigma_l['G'])
17             logW_l = sigma_l['logW']
18             Glist.append(G_l)
19             logWlist.append(logW_l)
20
21             weighted_samples.append((r_l,logW_l))
22
23             ELBO = sum(logWlist)
24
25             ghat = elbo_grad(Glist,logWlist)
26
27             sigma['Q'], max_norm = optimizer_step( sigma['Q'],ghat,t) #update
the proposal
28             print("results on iteration {} are ".format(t), sigma['Q'])
29             print("the max gradient is ", max_norm )
30
31
32
33             print("sigma['Q'] is ",sigma['Q'])
34
35             return weighted_samples, sigma['Q']
36

```

Sadly my importance sampling just was not working for some reason :(

2 Program 1


```
returning ghat: {'sample2': array([-0.49577966, -0.0261793 ], dtype=float32)}
norm of gradient: 0.49647036
results on iteration 93 are {'sample2': Normal(loc: 7.7331366539001465, scale: 0.91798853874
20654)}
the max gradient is 0.49647036
ELBO is tensor(-951.8278, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.5220273 , 0.08419174], dtype=float32)}
norm of gradient: 0.5287729
results on iteration 94 are {'sample2': Normal(loc: 7.728116989135742, scale: 0.917835891246
7957)}
the max gradient is 0.5287729
ELBO is tensor(-968.4595, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.42646474, -0.03271492], dtype=float32)}
norm of gradient: 0.42771772
results on iteration 95 are {'sample2': Normal(loc: 7.724055290222168, scale: 0.918322205543
5181)}
the max gradient is 0.42771772
ELBO is tensor(-951.6724, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.39113718, 0.0641632 ], dtype=float32)}
norm of gradient: 0.39636502
results on iteration 96 are {'sample2': Normal(loc: 7.720365524291992, scale: 0.918134987354
2786)}
the max gradient is 0.39636502
ELBO is tensor(-975.7202, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.42080986, 0.08003196], dtype=float32)}
norm of gradient: 0.4283527
results on iteration 97 are {'sample2': Normal(loc: 7.716432571411133, scale: 0.918498694896
698)}
the max gradient is 0.4283527
ELBO is tensor(-1007.1896, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.38780287, -0.11757746], dtype=float32)}
norm of gradient: 0.40523517
results on iteration 98 are {'sample2': Normal(loc: 7.712841987609863, scale: 0.918948173522
9492)}
the max gradient is 0.40523517
ELBO is tensor(-1002.4915, grad_fn=<AddBackward0>)
returning ghat: {'sample2': array([-0.44282976, -0.13143373], dtype=float32)}
norm of gradient: 0.46192318
results on iteration 99 are {'sample2': Normal(loc: 7.708779335021973, scale: 0.918293952941
8945)}
the max gradient is 0.46192318
sigma['Q'] is {'sample2': Normal(loc: 7.708779335021973, scale: 0.9182939529418945)}
{'sample2': Normal(loc: 7.708779335021973, scale: 0.9182939529418945)}
```

```

norm of gradient: 0.29127353
results on iteration 92 are {'sample2': Normal(loc: 7.540454864501953, scale: 0.915859937667
8467))}
the max gradient is 0.29127353
ELBO is tensor(-1230.8153, grad_fn=<AddBackward0>)
norm of gradient: 0.31071016
results on iteration 93 are {'sample2': Normal(loc: 7.537442684173584, scale: 0.915426969528
1982))}
the max gradient is 0.31071016
ELBO is tensor(-1245.2018, grad_fn=<AddBackward0>)
norm of gradient: 0.29013023
results on iteration 94 are {'sample2': Normal(loc: 7.534736156463623, scale: 0.915524959564
209))}
the max gradient is 0.29013023
ELBO is tensor(-1246.6462, grad_fn=<AddBackward0>)
norm of gradient: 0.3307241
results on iteration 95 are {'sample2': Normal(loc: 7.531586647033691, scale: 0.915119588375
0916))}
the max gradient is 0.3307241
ELBO is tensor(-1232.8281, grad_fn=<AddBackward0>)
norm of gradient: 0.26369676
results on iteration 96 are {'sample2': Normal(loc: 7.529160499572754, scale: 0.915133357048
0347))}
the max gradient is 0.26369676
ELBO is tensor(-1250.5586, grad_fn=<AddBackward0>)
norm of gradient: 0.3192326
results on iteration 97 are {'sample2': Normal(loc: 7.526177883148193, scale: 0.915463030338
2874))}
the max gradient is 0.3192326
ELBO is tensor(-1254.6959, grad_fn=<AddBackward0>)
norm of gradient: 0.26852316
results on iteration 98 are {'sample2': Normal(loc: 7.523730754852295, scale: 0.915418386459
3506))}
the max gradient is 0.26852316
ELBO is tensor(-1268.3943, grad_fn=<AddBackward0>)
norm of gradient: 0.26799098
results on iteration 99 are {'sample2': Normal(loc: 7.521484375, scale: 0.915682315826416)}
the max gradient is 0.26799098
sigma['Q'] is {'sample2': Normal(loc: 7.521484375, scale: 0.915682315826416)}
proposal is: {'sample2': Normal(loc: 7.521484375, scale: 0.915682315826416)}
Mean: tensor(0.0710, grad_fn=<SumBackward0>) Variance: tensor(3.3861, grad_fn=<SubBackward0
>)

```

But the sampled values look wrong:



3 Program2

Again we can see that the optimized proposal is looking correct.

```

the max gradient is 103.12931
ELBO is tensor(7425.1748, grad_fn=<AddBackward0>)
norm of gradient: 50.927845
norm of gradient: 70.857544
results on iteration 195 are {'sample1': Normal(loc: 2.3805956840515137, scale:
  0.1035614013671875), 'sample2': Normal(loc: 0.6931635737419128, scale: 8.503091
  812133789)}
the max gradient is 70.857544
ELBO is tensor(13869.1396, grad_fn=<AddBackward0>)
norm of gradient: 117.126724
norm of gradient: 436.69858
results on iteration 196 are {'sample1': Normal(loc: 2.3758506774902344, scale:
  0.10381706058979034), 'sample2': Normal(loc: 0.6936087608337402, scale: 8.50076
  5800476074)}
the max gradient is 436.69858
ELBO is tensor(1138.0188, grad_fn=<AddBackward0>)
norm of gradient: 34.049946
norm of gradient: 58.89527
results on iteration 197 are {'sample1': Normal(loc: 2.3748931884765625, scale:
  0.10417088121175766), 'sample2': Normal(loc: 0.6954396367073059, scale: 8.52292
  537689209)}
the max gradient is 58.89527
ELBO is tensor(11153.1797, grad_fn=<AddBackward0>)
norm of gradient: 215.83969
norm of gradient: 197.07481
results on iteration 198 are {'sample1': Normal(loc: 2.36855149269104, scale: 0
  .10431227833032608), 'sample2': Normal(loc: 0.6991506814956665, scale: 8.5205812
  45422363)}
the max gradient is 215.83969
ELBO is tensor(-1818.7682, grad_fn=<AddBackward0>)
norm of gradient: 175.42859
norm of gradient: 46.918015
results on iteration 199 are {'sample1': Normal(loc: 2.36621356010437, scale: 0
  .10518734157085419), 'sample2': Normal(loc: 0.7014965415000916, scale: 8.5297613
  1439209)}
the max gradient is 175.42859
sigma['Q'] is {'sample1': Normal(loc: 2.36621356010437, scale: 0.10518734157085
  419), 'sample2': Normal(loc: 0.7014965415000916, scale: 8.52976131439209)}
proposal is: {'sample1': Normal(loc: 2.36621356010437, scale: 0.1051873415708541
  9), 'sample2': Normal(loc: 0.7014965415000916, scale: 8.52976131439209)}

```

4 Program3

I have a bug that only affects program 3.

5 Program4

```
. 1.0000000000000000, 'sample1': Normal(loc: 0.1120022173700000, scale: 0.7070700000000000, 'sample10': Normal(loc: 0.028813054785132408, scale: 0.910309910774231), 'sample73': Normal(loc: -0.1274077445268631, scale: 1.211338043213), 'sample104': Normal(loc: 0.13367031514644623, scale: 1.000471591949463), 'sample9': Normal(loc: 0.23851779103279114, scale: 0.863947868347168), 'sample121': Normal(loc: 0.12602509558200836, scale: 0.90698998865), 'sample102': Normal(loc: -0.020959576591849327, scale: 1.05450439453125), 'sample55': Normal(loc: -0.716641902924, scale: 0.8955751657485962), 'sample87': Normal(loc: -0.0016292359214276075, scale: 0.82132261985), 'sample8': Normal(loc: 0.15166804194450378, scale: 0.9865884184837341), 'sample51': Normal(loc: 0.028155573826, scale: 0.897223174571991), 'sample72': Normal(loc: 0.28447407484054565, scale: 1.041612148284912), 'sample39': Normal(loc: -0.23673585057258606, scale: 0.9712879657745361), 'sample7': Normal(loc: 0.26283746957778, scale: 0.8800844550132751), 'sample54': Normal(loc: -0.008857478387653828, scale: 0.9268837571144104), 'sample101': Normal(loc: 0.14971597492694855, scale: 0.993392825126648), 'sample101': Normal(loc: 0.34996819496154785, scale: 0.9475797414779663), 'sample61': Normal(loc: -0.03995126858353615, scale: 0.9674408435821533), 'sample115': Normal(loc: -0.12889261543750763, scale: 1.0636239051818848), 'sample17': Normal(loc: -0.14979583024978638, scale: 0.9294111132621765), 'sample131': Normal(loc: 0.3166612386703491, scale: 1.0845746994018555), 'sample107': Normal(loc: 0.15554308891296387, scale: 1.019728422164917), 'sample109': Normal(loc: 0.09929069131612778, scale: 0.23301315308), 'sample25': Normal(loc: -0.2017146348953247, scale: 1.004259467124939), 'sample98': Normal(loc: 0.3766495883464813, scale: 0.9975871443748474), 'sample23': Normal(loc: -0.3070654273033142, scale: 0.98941212235), 'sample96': Normal(loc: 0.16982457041740417, scale: 0.9840357899665833), 'sample83': Normal(loc: 0.341820144653, scale: 0.9283028841018677), 'sample28': Normal(loc: -0.28709498047828674, scale: 0.87845247983932), 'sample117': Normal(loc: 0.12659983336925507, scale: 1.0993764400482178), 'sample27': Normal(loc: -0.37188762634, scale: 1.1278619766235352), 'sample88': Normal(loc: -0.3917408585548401, scale: 1.0515893697738647), 'sample116': Normal(loc: 0.19242410361766815, scale: 1.1206859350204468), 'sample5': Normal(loc: 0.3104699552059173, scale: 0.896874725818634), 'sample52': Normal(loc: -0.09707590937614441, scale: 1.0181375741958618), 'sample64': Normal(loc: 0.12425005435943604, scale: 1.166978120803833), 'sample46': Normal(loc: -0.05304451286792755, scale: 0.9874014258384705), 'sample129': Normal(loc: 0.02255052514374256, scale: 1.0709550380706787), 'sample57': Normal(loc: 0.07677019387483597, scale: 0.9037127494812012), 'sample53': Normal(loc: 0.3087914288043976, scale: 1.02441467285), 'sample12': Normal(loc: 0.1744071990251541, scale: 0.917870819568634), 'sample126': Normal(loc: 0.423068091273308, scale: 1.1834315061569214), 'sample93': Normal(loc: -0.03997012972831726, scale: 0.961156426), 'sample24': Normal(loc: 0.13767124712467194, scale: 1.0464117527008057), 'sample65': Normal(loc: 0.088283369217, scale: 0.9979518055915833), 'sample138': Normal(loc: -0.009301628917455673, scale: 0.93537604808807), 'sample29': Normal(loc: 0.22290989756584167, scale: 1.03630530834198), 'sample34': Normal(loc: -0.182759553102, scale: 0.9403443932533264), 'sample133': Normal(loc: 0.14275957643985748, scale: 1.0136116743087769), 'sample119': Normal(loc: -0.4713114798069, scale: 1.0176291465759277), 'sample86': Normal(loc: 0.1443261206150055, scale: 1.0197396278381348), 'sample136': Normal(loc: -0.03695037588477135, scale: 1.129204273223877), 'sample58': Normal(loc: 0.21542195975780487, scale: 0.9976047277450562), 'sample127': Normal(loc: 0.22007790207862854, scale: 0.459568858146667), 'sample66': Normal(loc: -0.48271283507347107, scale: 1.0355836153030396), 'sample95': Normal(loc: -0.29609158635139465, scale: 1.010400652885437), 'sample56': Normal(loc: 0.053470637649297714, scale: 0.97288894653), 'sample92': Normal(loc: 0.16050845384597778, scale: 1.168870449066162), 'sample48': Normal(loc: 0.424364298582077, scale: 1.1600818634033203), 'sample111': Normal(loc: -0.27430272102355957, scale: 1.16690818667), 'sample10': Normal(loc: 0.4499799907207489, scale: 0.9570946097373962)}  
the max gradient is 35.292686
```


6 Program5

7 Program3

7.1 results